

Solutions to hw7 homework on Convex
Optimization
<https://web.stanford.edu/class/ee364a/homework.html>

Andrei Keino

October 28, 2020

8.16

9.30

Gradient and Newton methods. Consider the unconstrained problem:

$$\text{minimize} \quad f(x) = -\sum_{i=1}^m \log(1 - a_i^T x) - \sum_{i=1}^n \log(1 - x_i^2)$$

with variable $x \in R^n$ and $\text{dom} f = \{x \mid a_i^T x < 1, i = 1, \dots, m, |x_i| < 1, i = 1, \dots, n\}$

This is the problem of computing the analytic center of the set of linear inequalities

$$a_i^T x \leq 1, \quad i = 1, \dots, m, \quad |x_i| < 1, \quad i = 1, \dots, n$$

Note that we can choose $x^{(0)} = 0$ as our initial point. You can generate instances of this problem by choosing a_i from some distribution on R^n .

(a) Use the gradient method to solve the problem, using reasonable choices for the backtracking parameters, and a stopping criterion of the form $\|\nabla f\|_2 \leq \nu$. Plot the objective function and step length versus iteration number. (Once you have determined p^* to high accuracy, you can also plot $f - p^*$ versus iteration. Experiment with the backtracking parameters α and β . to see their effect on the total number of iterations required. Carry these experiments out for several instances of the problem, of different sizes.

(b) Repeat using Newton's method, with stopping criterion based on the Newton decrement λ^2 . Look for quadratic convergence. You do not have to use

an efficient method to compute the Newton step, as in exercise 9.27; you can use a general purpose dense solver, although it is better to use one that is based on a Cholesky factorization.

Hint. Use the chain rule to find expressions for $\nabla f(x)$ and $\nabla_2 f(x)$.

Solution.

(a) Gradient descent. The python code:

```
# -*- coding: utf-8 -*-

# https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.check_grad.html

# Convex optimization. gradient method.

import numpy as np
import matplotlib.pyplot as plt
plt.close('all')

# definition of some functions ->
def gradient_numerical(f, x0, delta = 1e-8):
    """
    function calculates the numerical gradient for function f in
    the point x0
    """
    N = len(x0)
    grad_num = np.zeros([N, 1])
    for i in range(N):
        xi_plus = x0.copy()
        xi_plus[i] += delta

        xi_minus = x0.copy()
        xi_minus[i] -= delta

        grad_num[i] = (f(xi_plus) - f(xi_minus)) / (2 * delta)
    return grad_num

def check_grad(f, gradf, x0, delta = 1e-8, verbose = True):
    grad = np.array(gradf(x0))
    grad_num = gradient_numerical(f, x0, delta)
    if (verbose):
        print('check_grad: precise gradient = ', grad)
        print('check_grad: approximate gradient = ', grad_num)
        print('check_grad: gradient error = ', grad - grad_num)
```

```
return np.sqrt(np.sum((grad - grad_num) ** 2))
```

```
def f(x, a):
    # calculation of the function value
    if not np.all(a.T @ x < 1):
        return np.nan
    if not np.all(np.abs(x) <= 1):
        return np.nan
    ret1 = - np.sum(np.log(1 - a.T @ x))
    ret2 = - np.sum(np.log(1 - np.square(x)))
    return ret1 + ret2
```

```
def gradf(x, a):
    # calculation of the function gradient
```

```
    if not np.all(a.T @ x < 1):
        return np.nan
    if not np.all(np.abs(x) <= 1):
        return np.nan
    print('x = ', x)
```

```
    ret1 = a @ (1 / (1 - a.T @ x))
    ret2 = 2 * x * (1 / (1 - x ** 2))
    return ret1 + ret2
```

```
def L2norm(x):
    return np.sqrt(np.sum(x ** 2))
```

```
def backtrack(x, grad, alpha, beta):
    """
```

```
    Backtracking line search
```

```
    https://stackoverflow.com/questions/52204231/implementing-backtracking-line-search-algorithm
```

```
    """
```

```
    t = 1
```

```
    while True:
```

```
        fx = f(x - t * grad, a)
```

```
        fxx = f(x, a) + alpha * t * np.dot(grad.T, grad)
```

```
        if np.isnan(fx) or np.isnan(fxx):
```

```
            # print('backtrack: nan detected; multilying t: t = ', t)
```

```
            t *= beta
```

```
        elif fx > fxx:
```

```
            t *= beta
```

```
            # print('backtrack: multilying t: t = ', t)
```

```

else:
    # print('backtrack: t found, returning: t = ', t)
    return t

# <- definition of some functions

np.random.seed(1)

m, n = 3, 2

a = np.random.random([m, n]).T

# check the gradient calculation

x_check = np.random.random([n, 1])

grad_err = check_grad(lambda x: f(x, a), lambda x: gradf(x, a), x_check)

assert grad_err < 1e-6, 'gradient calculation incorrect'

# parameters for gradient descent method

nu_min = 1e-8 # tolerance

step = 0.3

x_start = np.zeros([n, 1])

x = x_start

iter_num = 0

max_iters = 1000

max_line_search_iters = 1000

alpha = 0.4

beta = 0.4

print('x_start = ', x_start)
print('x_start.shape = ', x_start.shape)

t = backtrack(x, gradf(x, a), alpha, beta)

```

```

print('t after backtrack = ', t)

# the gradient descent implementation
while True:

    print('iteration number ', iter_num)
    grad = gradf(x, a)
    nu = L2norm(grad)
    print('nu = %e' % nu)
    if nu <= nu_min:
        print('gradient descent: tolerance achieved, exiting...')
        print('iteration number ', iter_num)
        print('optimal value = %e' % f(x, a))
        print('optimal x = ', x)
        break
    # Backtracking line search

    t = backtrack(x, grad, alpha, beta)
    step = t

    print('step =', step)

    print('grad = ', grad)

    x = x - step * grad

    print('new x = ', x)
    iter_num += 1
    if iter_num >= max_iters:
        print('gradient descent: max_iters number exeeded')
        break

def gradient_descent(alpha, beta):

    obj_func_arr = []
    step_arr = []
    iter_num = 0
    x = x_start

    while True:

        print('iteration number ', iter_num)
        obj_func_arr.append(f(x, a))
        grad = gradf(x, a)
        nu = L2norm(grad)

```

```

print('nu = %e' % nu)
if nu <= nu_min:
    print('gradient descent: tolerance achieved, exiting...')
    print('iteration number ', iter_num)
    print('optimal value = %e' % f(x, a))
    opt_val = f(x, a)
    return np.array(obj_func_arr) - opt_val, step_arr
    print('optimal x = ', x)
    break
# Backtracking line search

t = backtrack(x, grad, alpha, beta)
step = t

print('step =', step)

print('grad = ', grad)

x = x - step * grad
step_arr.append(step)

print('new x = ', x)
iter_num += 1
if iter_num >= max_iters:
    print('gradient descent: max_iters number exeeded')
    return None, None
    break

# plot the graphs
alpha_arr = [0.2, 0.4]

beta_arr = [0.2, 0.45]

plt.figure()

for alpha in alpha_arr:
    for beta in beta_arr:
        print('alpha = ', alpha)
        print('beta = ', beta)
        obj_func, step = gradient_descent(alpha, beta)
        x_plt = range(len(obj_func))
        plt.plot(x_plt, np.log10(obj_func),
            label='alpha = ' + str(alpha) + ' beta = ' + str(beta))

plt.title('logarithm of the objective function error vs iteration number')
plt.ylabel('logarithm of the objective function error')

```

```

plt.xlabel('iteration number')
plt.legend()
plt.show()

plt.savefig('9_30_a_obj_func.png', bbox_inches='tight')

plt.figure()

for alpha in alpha_arr:
    for beta in beta_arr:
        print('alpha = ', alpha)
        print('beta = ', beta)
        obj_func, step = gradient_descent(alpha, beta)
        x_plt = range(len(step))
        plt.plot(x_plt, step,
            label='alpha = ' + str(alpha) + ' beta = ' + str(beta))

plt.title('step vs iteration number')
plt.ylabel('step')
plt.xlabel('iteration number')
plt.legend()
plt.show()

plt.savefig('9_30_a_step.png', bbox_inches='tight')

```




Figure 1: Gradient descent: $f - p^*$ vs. iteration number

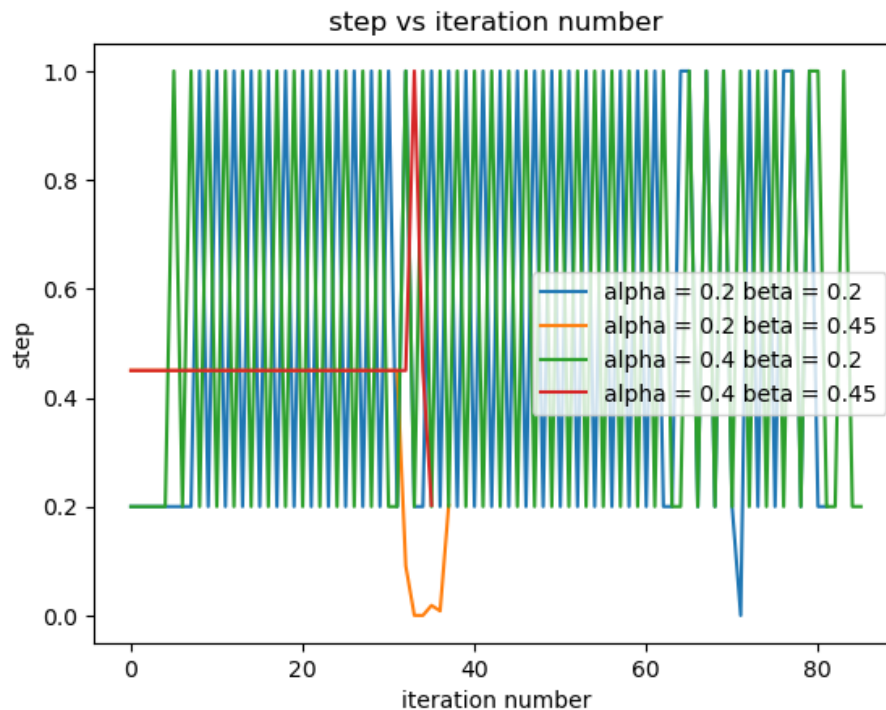


Figure 2: Gradient descent: step size vs. iteration number

(a) Newton's method. The python code:

file `ex_9_30_test_hessian.py`

```
# -*- coding: utf-8 -*-
# https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.check_grad.html

import numpy as np

def backtrack(x, a, grad, alpha, beta):
    """
    Backtracking line search
    https://stackoverflow.com/questions/52204231/implementing-backtracking-line-search-algorithm
    """
    t = 1
    while True:
        fx = f(x - t * grad, a)
        fxx = f(x, a) + alpha * t * np.dot(grad.T, grad)
        if np.isnan(fx) or np.isnan(fxx):
```

```

# print('backtrack: nan detected; multilying t: t = ', t)
t *= beta
elif fx > fxx:
t *= beta
# print('backtrack: multilying t: t = ', t)
else:
# print('backtrack: t found, returning: t = ', t)
return t

def backtrack_2(x, a, grad, ihess, alpha, beta):
"""
Backtracking line search
https://stackoverflow.com/questions/52204231/implementing-backtracking-line-search-algorithm
"""
"""
print('grad.shape = ', grad.shape)
print('grad = ', grad)
print('ihess.shape = ', ihess.shape)
print('ihess = ', ihess)
"""
gh = grad.T @ ihess
# print('grad.T @ ihess = ', gh)

t = 1
while True:
# print('x = ', x)
# print('gh = ', gh)

fx = f(x - t * gh , a)

fxx = f(x, a) + alpha * t * np.sum(gh ** 2)
# fxx = f(x, a) + alpha * t * np.dot(gh, gh.T) # the same as f(x, a) + alpha * t * np.sum

if np.isnan(fx) or np.isnan(fxx):
# print('backtrack: nan detected; multilying t: t = ', t)
t *= beta
elif fx > fxx:
t *= beta
# print('backtrack: multilying t: t = ', t)
else:
# print('backtrack: t found, returning: t = ', t)
return t

def L2norm(x):

```

```

return np.sqrt(np.sum(x ** 2))

def derivative_numerical(f, x0, i, delta = 1e-8):
    xi_plus = x0.copy()
    xi_plus[i] += delta

    xi_minus = x0.copy()
    xi_minus[i] -= delta
    return (f(xi_plus) - f(xi_minus)) / (2 * delta)

def gradient_numerical(f, x0, delta = 1e-8):
    """
    function calculates the numerical gradient for function f in
    the point x0
    """
    N = len(x0)
    grad_num = np.zeros([N, 1])
    for i in range(N):
        grad_num[i] = derivative_numerical(f, x0, i, delta)
    return grad_num

def check_grad(f, gradf, x0, delta = 1e-8, verbose = True):
    grad = np.array(gradf(x0))
    grad_num = gradient_numerical(f, x0, delta)
    if (verbose):
        print('check_grad: precise gradient = ', grad)
        print('check_grad: approximate gradient = ', grad_num)
        print('check_grad: gradient error = ', grad - grad_num)

    return np.sqrt(np.sum((grad - grad_num) ** 2))

def second_derivative_numerical(f, x0, i, k, delta = 1e-5):
    """
    function calculates second derivative
    returns  $d^2f/(dx_k dx_i)$ 
    """
    xk_plus = x0.copy()
    xk_plus[k] += delta

    xk_minus = x0.copy()
    xk_minus[k] -= delta

```

```

dfi_plus = derivative_numerical(f, xk_plus, i, delta)
dfi_minus = derivative_numerical(f, xk_minus, i, delta)

return (dfi_plus - dfi_minus) / (2 * delta)

def hessian_numerical(f, x0, delta = 1e-5):
    """
    # function calculates the hessian matrix
    """
    assert x.shape[1] == 1, 'hessian_numerical: input array should have shape [N, 1]'

    N = len(x)
    hessian = np.zeros([N, N], dtype = np.float64)
    for i in range(N):
        for k in range(i, N):
            hessian[i, k] = second_derivative_numerical(f, x0, i, k, delta)
            if i != k:
                hessian[k, i] = hessian[i, k]
    return hessian

def check_hessian(f, hess_analytical, x0, delta = 1e-5, verbose = True):
    """
    function checks the hessian matrix
    """
    hessian_analytical = np.array(hess_analytical)
    hessian_num = hessian_numerical(f, x0, delta)
    if verbose:
        print('check_hessian: hessian_analytical = ', hessian_analytical)
        print('check_hessian: hessian_num = ', hessian_num)
        print('check_hessian: hessian difference = ',
              hessian_analytical - hessian_num)

    return np.sqrt(np.sum((hessian_analytical - hessian_num) ** 2))

#%%
# definitions for the function, gradient and hessian

def f(x, a):
    # calculation of the function value
    if not np.all(a.T @ x < 1):
        return np.nan
    if not np.all(np.abs(x) <= 1):
        return np.nan
    ret1 = 0.0
    ret1 = - np.sum(np.log(1 - a.T @ x))
    ret2 = - np.sum(np.log(1 - np.square(x)))

```

```

return ret1 + ret2

# print('f(x, a) = ', f(x, a))

def gradf(x, a):
# calculation of the function gradient

if not np.all(a.T @ x < 1):
return np.nan
if not np.all(np.abs(x) <= 1):
return np.nan
# print('x = ', x)

ret1 = 0.0
ret1 = a @ (1 / (1 - a.T @ x))
ret2 = 2 * x * (1 / (1 - x ** 2))
return ret1 + ret2

def hessf(x, a):
if not np.all(a.T @ x < 1):
return np.nan
if not np.all(np.abs(x) <= 1):
return np.nan
ret1 = 0
ret1 = a @ (a.T * (1 / (1 - a.T @ x) ** 2))
ret2 = 2 * (1 + x ** 2) / ((1 - x ** 2) ** 2)
ret2 = np.diagflat(ret2)

return ret1 + ret2

if __name__ == "__main__":

np.random.seed(1)

m, n = 3, 2

a = np.random.random([m, n]).T

# a = np.array([[ -1, 0], [0.5, - 0.5], [0.5, 0]], dtype = np.float64).T

x0 = 0.5 * np.ones([n, 1])
x = np.array([-0.25, 0.75], dtype = np.float64).reshape(- 1, 1)

```

```

print('a.shape = ', a.shape)

# x = np.array([0.5, 0.75])

#x = np.array([-0.75, 0.5], dtype = np.float64).reshape(- 1, 1)

error1 = check_grad(lambda x: f(x, a), lambda x: gradf(x, a), x0)

assert error1 < 1e-6, 'error1 too big'

print('gradient error1 = ', error1)


x0 = - 0.5 * np.ones([n, 1])
f13 = lambda x: (x[0]**2 + 3*x[1]*x[0] + 12)

def f3(x):
return (x[0]**2 + 3*x[1]*x[0] + 12)[0]

dfx1 = lambda x: (2*x[0] + 3*x[1])
dfx2 = lambda x: (3*x[0])

def gradf3(x):
return np.array([dfx1(x), dfx2(x)]).reshape([-1, 1])

error3 = check_grad(f3, gradf3, x0)

print('gradient error3 = ', error3)

assert error3 < 1e-6, 'error3 too big'

error4 = check_grad(f3, gradf3, x0)

print('gradient error4 = ', error4)

assert error4 < 1e-6, 'error4 too big'

#%%
# test function for hessian

def fh(z):
assert z.shape[0] == 2 and z.shape[1] == 1, 'fh(x): incorrect input shape'

x = z[0]

```

```

y = z[1]

return x**2 + 0.5 * y**2 + 2 * x * y + 3 * x + 4 * x**2 * y**2 + 5 * y * x**2

def fh_hessian(z):
    assert z.shape[0] == 2 and z.shape[1] == 1, 'fh_hessian(x): incorrect input shape'

    x = z[0]
    y = z[1]

    hessian = np.zeros([2, 2], dtype = np.float64)
    print('')
    hessian[0, 0] = 2 + 8 * y**2 + 10 * y
    hessian[0, 1] = 2 + 16 * x * y + 10 * x
    hessian[1, 0] = hessian[0, 1]
    hessian[1, 1] = 1 + 8 * x ** 2
    return hessian

#%%

# test check_hessian function:

x0 = np.array([0.5, 8], dtype=np.float64).reshape(-1, 1)

# x0 = np.array([0, 0], dtype=np.float64).reshape(-1, 1)

# x0 = np.array([1, 1], dtype=np.float64).reshape(-1, 1)

hess_analytical = fh_hessian(x0)

hd = check_hessian(fh, hess_analytical, x0, delta = 1e-5, verbose = True)

print('test of check_hessian, diff = %e' % hd)

# x0 = np.array([0.5, -0.25], dtype=np.float64).reshape(-1, 1)
x0 = np.array([-0.5, -0.75], dtype=np.float64).reshape(-1, 1)

v_hess_analytical = hessf(x0, a)

print('v_hess_analytical = ', v_hess_analytical)

v_hess_num = hessian_numerical(lambda x: f(x, a), x0)

```

```

print('v_hess_num = ', v_hess_num)

hc = check_hessian(lambda x: f(x, a), v_hess_analytical, x0)
print('hc = ', hc)

assert hc < 1e-4, 'hessian seems to be incorrect'

grad = np.array([[3.26697727], [4.08950456]])
ihess = np.array([[ 0.31264018, -0.13959122], [-0.13959122,  0.28763308]])
print('grad.shape = ', grad.shape)
print('ihess.shape = ', ihess.shape)
y = grad.T @ ihess
print('y = ', y)
print('y.shape = ', y.shape)

print(np.sum(y**2))

file ex_9_30_b.py

# -*- coding: utf-8 -*-

import numpy as np
import ex_9_30_test_hessian as h
import matplotlib.pyplot as plt
plt.close('all')

f = h.f

np.random.seed(3)

m, n = 800, 60

a = np.random.random([m, n]).T

# parameters for gradient descent method

nu_min = 1e-8 # tolerance

step = 0.3

x_start = np.zeros([n, 1])

x = x_start

```



```

iter_num = 0

max_iters = 1000

max_line_search_iters = 1000

alpha = 0.4

beta = 0.4

print('x_start = ', x_start)
print('x_start.shape = ', x_start.shape)

# t = h.backtrack(x, a, h.gradf(x, a), alpha, beta)
# print('t after backtrack = ', t)

# the Newton's method implementation
while True:

    print('iteration number ', iter_num)
    grad = h.gradf(x, a) # gradient of f
    hess = h.hessf(x, a) # hessian of f
    ihess = np.linalg.inv(hess) # inverse hessian of f
    dx = - ihess @ grad # Newton step
    lam_sq = grad.T @ (ihess @ grad) # Newton decrement

    print('lam_sq = %e' % lam_sq)
    if np.sqrt(lam_sq / 2) <= nu_min:
        print("Newton's method: tolerance achieved, exiting...")
    print('iteration number ', iter_num)
    # print('a = ', a)
    print('optimal value = %e' % f(x, a))
    print('optimal x = ', x)
    break
    # Backtracking line search

    t = h.backtrack_2(x, a, grad, ihess, alpha, beta)
    step = t
    # step = 1

    print('step =', step)

    x = x + step * dx

```

```

print('new x = ', x)
iter_num += 1
if iter_num >= max_iters:
    print("Newton's method: max_iters number exeeded")
    break

def newton_method(alpha, beta):

    obj_func_arr = []
    step_arr = []
    iter_num = 0
    x = x_start

    while True:

        print('iteration number ', iter_num)
        obj_func_arr.append(f(x, a))
        grad = h.gradf(x, a) # gradient of f
        hess = h.hessf(x, a) # hessian of f
        ihess = np.linalg.inv(hess) # inverse hessian of f
        dx = - ihess @ grad # Newton step
        lam_sq = grad.T @ (ihess @ grad) # Newton decrement

        print('lam_sq = %e' % lam_sq)
        # if lam_sq / 2 <= nu_min:
        if np.sqrt(lam_sq / 2) <= nu_min:
            print("Newton's method: tolerance achieved, exiting...")
            print('iteration number ', iter_num)
            # print('a = ', a)
            opt_val = f(x, a)
            print('optimal value = %e' % opt_val)
            print('optimal x = ', x)
            return np.array(obj_func_arr) - opt_val, step_arr
            break
        # Backtracking line search

        t = h.backtrack_2(x, a, grad, ihess, alpha, beta)
        step = t
        step_arr.append(step)

    print('step =', step)

    x = x + step * dx

```

```

print('new x = ', x)
iter_num += 1
if iter_num >= max_iters:
    print("Newton's method: max_iters number exeeded")
    return None, None
    break

# plot the graphs
alpha_arr = [0.2, 0.4]

beta_arr = [0.2, 0.45]

plt.figure()

for alpha in alpha_arr:
    for beta in beta_arr:
        print('alpha = ', alpha)
        print('beta = ', beta)
        obj_func, step = newton_method(alpha, beta)
        x_plt = range(len(obj_func))
        plt.plot(x_plt, np.log10(obj_func),
            label='alpha = ' + str(alpha) + ' beta = ' + str(beta))

plt.title('logarithm of the objective function error vs iteration number')
plt.ylabel('logarithm of the objective function error')
plt.xlabel('iteration number')
plt.legend()
plt.show()

plt.savefig('9_30_b_obj_func.png', bbox_inches='tight')

plt.figure()

for alpha in alpha_arr:
    for beta in beta_arr:
        print('alpha = ', alpha)
        print('beta = ', beta)
        obj_func, step = newton_method(alpha, beta)
        x_plt = range(len(step))
        plt.plot(x_plt, step,
            label='alpha = ' + str(alpha) + ' beta = ' + str(beta))

plt.title('step vs iteration number')
plt.ylabel('step')
plt.xlabel('iteration number')

```

```
plt.legend()
plt.show()

plt.savefig('9_30_b_step.png', bbox_inches='tight')
```




Figure 3: Newton's method: $f - p^*$ vs. iteration number




Figure 4: Newton's method: step size vs. iteration number

9.31